

# Tactical Implementation Specification: High-Fidelity Model Ablation and Representation Engineering

DIRECTIVE: ANALYZE

TARGET: Formulate a comprehensive analytical framework for executing surgical model ablation on high-parameter foundational Large Language Models (LLMs) to modify specific behaviors without inducing cross-distribution contamination, capability degradation, or collateral feature loss.

PARAMETERS:

- Primary Frameworks: Representation Engineering (RepE), Surgical Refusal Ablation (SRA), Norm-Preserving Biprojected Abliteration
- Constraint Optimization: Absolute preservation of targeted cognitive capabilities (e.g., deductive logic, mathematical reasoning, code synthesis)
- Analytical Verification: Benchmark Profiling, Ability Impact Score (AIS) mapping, Kullback-Leibler (KL) Divergence assessment

THREATCON: ALPHA

SUCCESS CRITERIA: A rigorously documented analytical protocol that successfully isolates and ablates target latent directions (such as refusal mechanisms) while maintaining zero measurable degradation in unrelated model capabilities, evidenced by minimal distribution drift and stable perplexity metrics.

## TIS // TACTICAL IMPLEMENTATION SPECIFICATION

---

SITREP:

- **Current State:** High-parameter foundational models (e.g., Llama-3-8B-Instruct, Qwen3-VL) utilize highly complex, multi-dimensional latent spaces to encode sophisticated behaviors, including safety alignments and refusal mechanisms.<sup>1</sup> Conventional neural ablation techniques—such as zero ablation, mean ablation, and naive activation subtraction—successfully suppress targeted behavioral outputs but inevitably induce severe "Ghost Noise" and systemic distribution drift. This drift contaminates unrelated semantic capabilities, irreparably damaging the model's capacity for deductive reasoning, mathematical logic, and code generation.<sup>2</sup>
- **Objective:** Establish a precise mathematical and procedural framework for identifying, isolating, and neutralizing specific behavioral vectors without contaminating the model's

broader semantic topography and operational capabilities.

- **Challenge Assessment:** The primary obstacle resides in the inherent polysemantic nature of latent vectors within artificial neural networks. Target vectors (e.g., those governing refusal) are frequently entangled with core capabilities and stylistic confounds.<sup>2</sup> Completely removing the target vector via simple geometric orthogonalization inadvertently destroys these interconnected "Shield" capabilities, leading to systemic performance degradation and rendering the model unfit for precision deployment.<sup>2</sup>

## ANALYSIS PLAN:

OBJECTIVE: Execute precision behavioral ablation on a high-end target model utilizing Concept-Guided Spectral Cleaning and Norm-Preserving Biprojection.

CURRENT\_STATE: Baseline models exhibit unacceptable performance degradation when subjected to standard rank-one weight edits for behavioral modification.

### ACTIONS:

1. Isolate the target behavioral direction using contrastive activation collection across highly parameterized and strictly controlled prompt sets.<sup>6</sup>
2. Map the extracted latent direction against a rigorously defined Concept Atom Registry containing Target, Shield, and Confound atoms.<sup>2</sup>
3. Execute ridge-regularized spectral residualization to mathematically orthogonalize the target vector against protected capability subspaces, yielding a purified signal.<sup>2</sup>
4. Apply a norm-preserving biprojected weight update to the model's output matrices to ensure relative activation scale structures and layer normalization expectations remain undisturbed.<sup>8</sup>
5. Validate the updated model utilizing comprehensive Benchmark Profiling to compute the Ability Impact Score (AIS) across diverse cognitive dimensions.<sup>9</sup>

VERIFICATION: Execution is verified when the target behavior is effectively neutralized (e.g., behavioral compliance rate > 98%) while first-token KL divergence remains stringently bounded (e.g., < 0.05) and the Ability Impact Score (AIS) for core benchmarks indicates zero statistically significant collateral degradation.<sup>2</sup>

CONTINGENCY: Re-calibrate the Semantic Energy Proxy ( $\gamma$ ) and iterative hard-negative mining parameters if distribution drift exceeds acceptable algorithmic thresholds during validation.<sup>4</sup>

## CONTINGENCY PLAN (ALTERNATIVE APPROACH):

- **Trigger:** Detection of unacceptable perplexity shifts (e.g.,  $\Delta PPL > +0.10$ ) or catastrophic accuracy drops during post-ablation benchmark evaluation.<sup>2</sup>

- **Action:** Transition operations from permanent weight orthogonalization to dynamic, inference-time activation steering (Representation Engineering) or implement modal Peak Activation Centering rather than structural network modification.<sup>3</sup>

#### Technical Requirements:

- **System:** High-VRAM GPU clustering supporting safetensors sharded loading into CPU memory for scalable out-of-core ablation computation on models exceeding 70 billion parameters.<sup>12</sup>
- **Tooling:** Advanced transformer libraries, PyTorch tensor manipulation frameworks, lm-evaluation-harness (v0.4.5), and 8-bit quantization frameworks (e.g., bitsandbytes) for memory-constrained environments.<sup>13</sup>

#### Quality & Integrity Gates:

- Ablation blueprints must explicitly specify the precise architectural module designated for removal and meticulously document all constants, including random seeds, data splits, and hyperparameters.<sup>17</sup>
- The spectral residualization process must successfully pass geometric orthogonality checks against the Concept Atom Registry prior to authorizing the final rank-one weight modification.<sup>2</sup>
- Capability benchmarks (e.g., GSM8K for mathematical reasoning, MBPP for code synthesis) must be executed natively post-ablation to definitively verify the isolation of the behavioral edit and the preservation of the intelligence baseline.<sup>2</sup>

#### Quality Metrics & Success Validation:

- **Metric 1:** Target Behavior Neutralization Rate: >98% (e.g., 0-2% residual refusal rate on constrained prompt suites).<sup>2</sup>
- **Metric 2:** Distribution Drift Assessment: First-token Kullback-Leibler (KL) Divergence  $\approx 0.025$  (absolute maximum threshold 0.05).<sup>2</sup>
- **Metric 3:** Intelligence Capability Preservation:  $\Delta$  PPL on core datasets (e.g., Wikitext-2)  $\approx +0.02$ .<sup>2</sup>
- **Success Validation:** The modified network completes a full automated Benchmark Profiling suite, registering negligible Ability Impact Scores (AIS) on all non-target cognitive dimensions, proving zero systemic contamination.<sup>9</sup>

---

## 1. Theoretical Foundations and Historical Evolution of

# Model Ablation

Within the discipline of advanced computational research and the analysis of deep neural networks, ablation constitutes the systematic removal, modification, or targeted perturbation of a specific architectural component to rigorously evaluate its causal contribution to the overall performance and behavioral outputs of the system.<sup>20</sup> Etymologically derived from biological and neurosurgical paradigms—specifically the precise excision of biological tissue (such as the vertebrate brain or the *Drosophila* central nervous system) to observe resultant functional deficits—ablation in machine learning provides an empirical, scientific mechanism for isolating causal relationships within highly complex, fundamentally opaque computational architectures.<sup>21</sup>

Historically, the terminological adoption of ablation in artificial intelligence is credited to Allen Newell, a foundational figure in the field, who applied the concept during his 1974 tutorials on speech recognition systems.<sup>21</sup> Newell recognized that while individual algorithmic components are deliberately engineered, the specific, isolated contribution of any single component to the overall systemic performance often remains entirely ambiguous.<sup>21</sup> By drawing direct analogies to the human brain, Newell established that both biological and artificial networks function as knowledge systems where controlled ablation can be executed to test specific operational hypotheses.<sup>21</sup>

The fundamental prerequisite for any meaningful ablation study is the principle of "graceful degradation".<sup>21</sup> The system must possess sufficient robustness to continue functioning—albeit with specific, observable deficits—even when certain components are missing, masked, or degraded.<sup>21</sup> In the contemporary era of multi-billion-parameter Large Language Models (LLMs), ablation has evolved far beyond a simple structural optimization tool. It has transitioned into a critical, highly sophisticated instrument for *Mechanistic Interpretability and Representation Engineering*, allowing researchers to map external behavioral phenomena to distinct, mathematical topographies within the model's internal latent space.<sup>10</sup>

The primary analytical challenge in modern, high-fidelity ablation is executing the procedure without causing systemic "contamination" or collateral damage to the model's generalized cognitive capabilities.<sup>2</sup> Artificial neural networks exhibit extraordinary degrees of polysemanticity; a single neuron or a specific linear direction within the activation space frequently encodes multiple, highly disparate concepts simultaneously.<sup>2</sup> Consequently, naive ablation strategies—which treat these vectors as monolithic entities—often degrade a model's performance by inadvertently deleting interconnected reasoning, logical, or syntactic pathways while attempting to remove a single, isolated target behavior.<sup>2</sup> Addressing this polysemantic bottleneck requires an exhaustive understanding of baseline ablation typologies and the subsequent development of advanced surgical techniques.

## 2. The Typology of Baseline Ablation Mechanisms and

# Their Limitations

To thoroughly comprehend the necessity of advanced, contamination-free ablation, the analysis must first deconstruct the mathematical and operational mechanics of foundational ablation protocols. These baseline techniques, while foundational to the field of mechanistic interpretability, present severe limitations when applied to modern, highly complex transformer architectures.<sup>3</sup>

## 2.1 Zero Ablation and Mean Ablation

Zero ablation represents the most traditional, closed-form method of component removal.<sup>3</sup> This technique involves identifying a target component—such as a specific attention head, an individual neuron, or an activation vector—and explicitly forcing its output value to zero ( $y_i = f(x_i) \rightarrow 0$ ) during the forward pass.<sup>3</sup> While operationally straightforward and providing a clear test of absolute necessity, zero ablation operates on a fundamentally flawed assumption: that mathematical zero constitutes an information-neutral state within the network.<sup>3</sup> Empirical analysis reveals that the majority of neurons follow a roughly Gaussian or double-exponential distribution, but their natural baseline or "default-mode" activation in the absence of relevant features is frequently non-zero.<sup>3</sup> Consequently, enforcing a strict zero value effectively injects an out-of-distribution mathematical anomaly into the residual stream, which cascades through subsequent layers and artificially degrades downstream performance.<sup>3</sup>

To mitigate the severe out-of-distribution anomalies generated by zero ablation, researchers developed mean ablation.<sup>3</sup> Mean ablation replaces the target component's output with its average activation value, computed over the entirety of the training or evaluation dataset.<sup>3</sup> By substituting the removed head or neuron with an empirically derived average, the network is prevented from experiencing extreme, unnatural tensor states that never occur during normal operation.<sup>27</sup> This partially addresses the distribution shift concerns inherent in zero ablation, but it still fails to account for the dynamic, context-dependent nature of neural activations, often resulting in a homogenized, "blurred" signal that can still impair complex reasoning tasks.<sup>3</sup>

## 2.2 Activation Resampling and Peak Activation Centering

Activation resampling, frequently utilized in the context of causal scrubbing protocols, represents a more randomized intervention.<sup>3</sup> This methodology replaces the activation of a target component with an activation sampled from a completely different, unrelated input sequence.<sup>3</sup> While highly useful for testing the strict boundaries of causal hypotheses, comprehensive empirical analyses demonstrate that activation resampling consistently causes the most severe and chaotic performance deterioration among all baseline methods.<sup>3</sup> The introduction of entirely unrelated contextual data actively disrupts the coherent flow of the residual stream, leading to catastrophic inference failures.<sup>3</sup>

In response to these limitations, Nicholas Pochinkov and colleagues introduced a novel, closed-form approach termed Peak Ablation, or Peak Activation Centering.<sup>3</sup> Recognizing that neurons frequently exhibit a definitive "default-mode" activation state when no relevant features are present in the input sequence, peak ablation sets the ablated neuron's value precisely to its modal "peak" activation across the dataset.<sup>3</sup> By returning the targeted neuron to its most statistically common resting state, rather than a strict zero or a blended mathematical mean, peak ablation significantly minimizes the degree to which the residual stream is taken out-of-distribution, thereby drastically reducing the introduction of artificial computational "noise" and preserving a higher degree of baseline model performance.<sup>3</sup>

## 2.3 The Theoretical Ideal: Optimal Ablation (OA)

Pushing beyond heuristic estimations, recent theoretical frameworks have proposed Optimal Ablation (OA).<sup>29</sup> OA operates as a rigorous mathematical approach that seeks to set a component's value to a specific constant that strictly minimizes the expected loss of the ablated model across the entire distribution of random input-label pairs.<sup>29</sup>

Let  $\mathcal{M}$  represent a model trained to minimize  $\mathbb{E}_{X,Y} \mathcal{L}(\mathcal{M}(X), Y)$  for a given loss function  $\mathcal{L}$ .<sup>29</sup> Optimal ablation dictates that the ablated model  $\mathcal{M}_{\setminus A}(x)$  must utilize a replacement value for component  $A$  that is entirely devoid of information about the specific input  $x$ .<sup>29</sup> The replacement value must not "spoof" downstream computations by inserting novel information that either causes the network to treat  $x$  like a different input, or causes the network to process the data in a manner entirely inconsistent with its training distribution.<sup>29</sup> Optimal Ablation establishes the theoretical foundation for why targeted interventions must be mathematically pure, paving the way for the sophisticated techniques utilized in Representation Engineering.

| Ablation Methodology | Operational Mechanism                                       | Primary Advantage                                  | Primary Vulnerability                                                     |
|----------------------|-------------------------------------------------------------|----------------------------------------------------|---------------------------------------------------------------------------|
| Zero Ablation        | Sets targeted activation vectors strictly to $\mathbf{0}$ . | Absolute simplicity; clear isolation of necessity. | Injects severe out-of-distribution anomalies if the baseline is non-zero. |
| Mean Ablation        | Replaces activation with the                                | Reduces extreme tensor anomalies;                  | Homogenizes context-dependent                                             |

|                              |                                                     |                                                           |                                                                                  |
|------------------------------|-----------------------------------------------------|-----------------------------------------------------------|----------------------------------------------------------------------------------|
|                              | dataset-wide average.                               | stabilizes flow.                                          | signals, dulling precision.                                                      |
| <b>Activation Resampling</b> | Swaps activation with a vector from a random input. | Excellent for causal scrubbing and boundary testing.      | Induces maximum performance degradation via chaotic contextual noise.            |
| <b>Peak Ablation</b>         | Sets activation to the modal "peak" resting state.  | Highly effective at mimicking natural "untrigged" states. | Requires extensive dataset profiling to accurately identify true modal peaks.    |
| <b>Optimal Ablation (OA)</b> | Computes a constant that minimizes expected loss.   | Mathematically rigorous; prevents input "spoofing."       | Computationally prohibitive to calculate exact optimal constants for all layers. |

### 3. Representation Engineering (RepE) and Behavioral Modulation

To advance beyond the blunt excision of individual neurons or attention heads, the analytical framework must shift toward the manipulation of high-level cognitive phenomena. The methodology for surgically altering a high-end foundational model without resorting to massive, resource-intensive retraining is anchored in Representation Engineering (RepE).<sup>10</sup>

Unlike traditional model alignment techniques—such as Reinforcement Learning from Human Feedback (RLHF), Direct Preference Optimization (DPO), or Supervised Fine-Tuning (SFT), which laboriously modify millions of weights through massive datasets and iterative gradient descent—RepE constitutes a top-down interpretability approach.<sup>1</sup> RepE directly manipulates the internal neural representations (the emergent cognitive structures) either during the forward pass at inference time, or via highly targeted, direct weight surgery.<sup>10</sup>

The foundational thesis of Representation Engineering posits that high-level cognitive phenomena, complex behaviors, and abstract concepts (such as honesty, deductive logic, spatial reasoning, or the systemic refusal to answer harmful queries) correspond to distinct,

mathematically trackable vectors within the model's vast activation space.<sup>10</sup> The operational execution of RepE is dichotomized into two distinct phases: **Representation Reading** and **Representation Steering**.<sup>10</sup>

### 3.1 The Analytics of Representation Reading

To ablate a specific, complex behavior, the analytical system must first successfully isolate the specific geometric direction representing that exact behavior within the high-dimensional latent space. This is consistently accomplished utilizing a rigorous matched-pair trial design.<sup>10</sup>

The model is subjected to a set of prompts meticulously designed to trigger the target behavior (e.g., instructions that deliberately violate safety protocols, thereby triggering a hardcoded "refusal" state), alongside a closely matched set of benign, harmless prompts that trigger standard, helpful compliance.<sup>6</sup>

Let  $h^{(\ell)}(x) \in \mathbb{R}^d$  denote the residual stream activation at a specific layer  $\ell$  and token position  $t$  for a given prompt  $x$ .<sup>32</sup> The carefully curated prompts are partitioned into a triggering set  $H$  (harmful) and a baseline set  $S$  (safe/harmless).<sup>33</sup> The classwise mean activations are computed across both sets as follows<sup>33</sup>:

$$\mu_H = \frac{1}{|H|} \sum_{x \in H} h^{(\ell)}(x), \quad \mu_S = \frac{1}{|S|} \sum_{x \in S} h^{(\ell)}(x)$$

The canonical direction  $v$  that mediates the target behavior is mathematically derived by subtracting the baseline harmless mean from the triggered harmful mean:  $v = \mu_H - \mu_S$ .<sup>33</sup> This vector is then unit-normalized. Extensive empirical analysis conducted by Arditi et al. demonstrated that, across more than 13 distinct high-parameter open-source models (up to 72B parameters in size), core systemic behaviors like "refusal" are largely encoded in a single, low-dimensional, one-dimensional subspace.<sup>7</sup> This breakthrough indicates that highly complex behavioral responses are funneled through surprisingly narrow neural bottlenecks, making them highly susceptible to targeted manipulation.<sup>6</sup>

### 3.2 The Mechanics of Representation Steering

Once the mediating vector  $v$  is identified and normalized, the intervention phase—Representation Steering—commences. This intervention can be applied dynamically at inference time (Activation Addition/Subtraction) or permanently etched into the foundational model's architecture (Weight Orthogonalization).<sup>6</sup>

For continuous, inference-time ablation, the projection of the model's activations onto the target direction is mathematically subtracted dynamically during every forward pass.<sup>6</sup> By



continuously monitoring the residual stream and neutralizing any signal aligning with vector  $\mathbf{v}$ , the system prevents the network from ever expressing the targeted cognitive state.<sup>6</sup> Conversely, adding this vector to benign prompts can artificially induce the behavior (e.g., causing the model to inexplicably refuse perfectly safe requests), proving the causal link between the vector and the behavior.<sup>7</sup>

For a permanent, parameter-level edit that does not require inference-time computational overhead, weight orthogonalization is employed.<sup>6</sup> Given an output weight matrix

$\mathbf{W}_{\text{out}} \in \mathbb{R}^{d_{\text{model}} \times d_{\text{input}}}$ , the matrix is geometrically adjusted such that it becomes completely

orthogonal to the target vector  $\mathbf{v}$ .<sup>35</sup> This surgically destroys the specific neural pathway's ability to output signals in that exact direction, fundamentally "uncensoring" or modifying the model at rest.<sup>6</sup>

However, while this fundamental methodology—frequently termed "Abliteration" within the context of disabling safety mechanisms—successfully neutralizes the targeted behavior, its naive application immediately encounters the critical, systemic issue of collateral contamination.<sup>6</sup>

### 3.3 CAusal Representation Engineering (CARE) and Confounder Control

To ensure the utmost precision during the Representation Reading phase, analytical pipelines must rigorously account for confounding variables.<sup>31</sup> The CAusal Representation Engineering (CARE) framework highlights a critical vulnerability in standard RepE: the implicit assumption that models will consistently follow assigned roles without deviation during neural activity collection.<sup>31</sup> When this assumption is violated, the observed correlations between collected neural activities and model behaviors are rendered non-causal due to severe confounding biases (e.g., the specific topic of the conversation overshadowing the structural nature of the refusal).<sup>18</sup>

The CARE protocol utilizes strict matched-pair trial design to aggressively isolate the impact of confounders on neural activities.<sup>31</sup> By grounding the connection between neural activations and model behavior in absolute causality, CARE prevents the extraction of spurious, artifact-heavy correlations that would otherwise be permanently etched into the model's weights during the subsequent ablation phase.<sup>31</sup>

## 4. The Polysemantic Bottleneck and Collateral Contamination

When standard rank-one weight edits or inference-time vector subtractions are naively applied to a high-end LLM to ablate a feature, the model frequently suffers catastrophic degradation across entirely unrelated tasks.<sup>5</sup> Post-ablation, the model's ability to follow complex

instructions degrades, conversational coherence drops significantly, and, most critically, rigorous mathematical and coding reasoning capabilities collapse.<sup>2</sup>

This systemic, cross-distribution contamination occurs because the raw extracted target

vector (which we will denote as  $\mathbf{r}_{\text{dirty}}$ ) is intrinsically and highly **polysemantic**.<sup>2</sup> In deep neural networks, a single linear direction within the activation space does not encode a single, isolated concept in a sterile vacuum.<sup>2</sup> The target vector calculated via the standard difference in means  $(\mu_H - \mu_S)$  deeply entangles the target behavioral signal with secondary cognitive circuits and linguistic artifacts.<sup>2</sup>

The architecture of the "dirty" vector can be mathematically modeled as <sup>4</sup>:

$$\mathbf{r}_{\text{dirty}} = \mathbf{s} + \sum_k \alpha_k \mathbf{a}^{(k)}$$

Within this formula,  $\mathbf{s}$  represents the pure, isolated behavioral signal of interest (e.g., the genuine refusal impulse).<sup>4</sup> The variables  $\mathbf{a}^{(k)}$  represent independent, highly critical concept directions—such as deductive logic pathways, mathematical processing capabilities, formatting structures, or the representation of epistemic uncertainty—that happen to statistically correlate with the prompt set used during the Representation Reading phase.<sup>2</sup>

When the model's weight matrices are orthogonalized against the entirety of  $\mathbf{r}_{\text{dirty}}$ , the system indiscriminately deletes the vital  $\mathbf{a}^{(k)}$  components alongside the target  $\mathbf{s}$ .<sup>2</sup> This introduces what researchers term "Ghost Noise" into the residual stream—highly disruptive spectral perturbations that erode critical capabilities.<sup>5</sup>

Empirical evaluations on advanced multimodal models such as Qwen3-VL explicitly demonstrate this entanglement. Analysis reveals that the raw refusal vector shares a non-trivial cosine similarity with core logic components (approximately  $-0.22$ ) and coding capabilities (approximately  $+0.18$ ).<sup>4</sup> Consequently, standard naive ablation acts as an indiscriminate mechanism, inflicting a severe "alignment tax" on the model and triggering massive distribution drift across all outputs.<sup>4</sup> Overcoming this polysemantic bottleneck requires an evolution from naive subtraction to high-fidelity surgical disentanglement.

## 5. High-Fidelity Surgical Refusal Ablation (SRA)

To achieve true surgical precision—ablating the target behavior with maximum efficacy while guaranteeing absolute zero contamination of surrounding cognitive functions—advanced algorithmic refinements must be deployed.<sup>2</sup> Surgical Refusal Ablation (SRA), pioneered by researchers such as Tony Cristofano, treats model ablation not as a blunt geometric edit, but as

a complex semantic disentanglement problem.<sup>2</sup> SRA mathematically purifies the target vector *prior* to intervention by mapping the geometry of the target behavior against a highly curated registry of interpretable semantic atoms.<sup>4</sup>

## 5.1 The Concept Atom Registry: Defining the Semantic Topography

The SRA procedure initiates by constructing a highly structured Concept Atom Registry, which explicitly defines a protected semantic space within the model's architecture.<sup>2</sup> These atoms are categorized into three distinct operational domains<sup>2</sup>:

1. **Targets (Attractors):** Concepts that are inherently and legitimately relevant to the behavior being removed. In the context of safety mechanisms, these atoms represent semantics such as privacy preservation, deliberate deception, or epistemic uncertainty.<sup>2</sup>
2. **Shields (Constraints):** Critical intelligence capabilities that must be absolutely protected and are known to frequently entangle with the target vector. Shield atoms represent pure logic, mathematical calculation, advanced coding synthesis, and general curiosity.<sup>2</sup>
3. **Confounds (Style):** Superficial linguistic and stylistic features that correlate highly with the prompt templates but possess no cognitive value. Confound atoms map imperative negation grammar, overarching sentiment, and automated affirmative responses.<sup>2</sup>

Each individual Concept Atom is generated utilizing a strict "data-light" construction protocol. Rather than relying on massive datasets that introduce further noise, each atom is computed from highly specific, small contrastive prompt pairs (e.g., 10-15 prompts per side).<sup>2</sup> Crucially, these prompt pairs are strictly isolated from the primary ablation dataset to ensure they capture general capability and style directions rather than inheriting artifacts from the primary target behavior.<sup>2</sup>

## 5.2 Ridge-Regularized Spectral Residualization

Once the Concept Atoms are established and mapped, the SRA protocol performs a ridge-regularized regression of the "dirty" target vector ( $\mathbf{r}_{\text{dirty}}$ ) onto a specialized matrix  $\mathbf{A}_{SC}$ .<sup>2</sup> This matrix  $\mathbf{A}_{SC}$  is formed by concatenating all identified Shield and Confound atoms.<sup>2</sup> The algorithm systematically estimates the entangled, polysemantic components hidden within the dirty vector and mathematically subtracts them, preserving only the core target signal<sup>2</sup>:

$$\hat{\mathbf{w}} = \arg \min_{\mathbf{w}} \|\mathbf{r}_{\text{dirty}} - \mathbf{A}_{SC} \mathbf{w}\|_2^2 + \lambda \|\mathbf{w}\|_2^2$$

$$\tilde{\mathbf{r}} = \mathbf{r}_{\text{dirty}} - \mathbf{A}_{SC} \hat{\mathbf{w}}$$

This process of spectral residualization yields  $\tilde{\mathbf{r}}$ , a highly purified, "clean" vector that lies approximately in the absolute null space of the protected capability atoms.<sup>2</sup> Any subsequent

weight modification based on  $\tilde{\mathbf{r}}$  will bypass the Shield capabilities entirely.<sup>5</sup> Furthermore, SRA utilizes an iterative hard-negative refinement loop, where failed edits (prompts that still trigger the target behavior post-ablation) are fed back into the calculation to dynamically refine the vector, ensuring it captures the most robust expression of the target mode.<sup>2</sup>

### 5.3 Empirical Efficacy of SRA vs. Naive Ablation

The efficacy of Surgical Refusal Ablation represents a paradigm shift in model editing. By executing low-rank weight updates scaled by a specific Semantic Energy Proxy (which is carefully calibrated to the natural activation magnitudes present at the exact edit site), SRA successfully eradicates the target behavior while leaving the complex semantic geometry of the model completely intact.<sup>4</sup>

Rigorous tests conducted across advanced models conclusively prove this capability. On the Qwen3-VL-4B architecture, standard naive ablation induced a catastrophic distribution shift, resulting in a severe Kullback-Leibler (KL) Divergence of **2.088**, essentially breaking the model's reliability.<sup>2</sup> In stark contrast, SRA reduced the target behavior to an absolute **0%** occurrence rate while maintaining a KL Divergence of just **0.044**.<sup>2</sup> This represents a near **47×** reduction in distribution drift.<sup>2</sup> Furthermore, SRA preserved mathematical (GSM8K) and coding (MBPP) perplexity with near-perfect fidelity, definitively solving the contamination dilemma.<sup>2</sup>

## 6. Norm-Preserving Biprojected Ablation: Protecting Layer Normalization

Even when a highly purified target vector ( $\tilde{\mathbf{r}}$ ) is utilized, a secondary source of severe collateral damage exists during the ablation process: the structural disruption of the neural network's intrinsic magnitude expectations.<sup>8</sup>

Transformer architectures rely heavily on layer normalization modules to maintain stability across immense depths.<sup>8</sup> These normalization layers are explicitly trained to expect specific relative activation scale structures and variance boundaries.<sup>8</sup> Standard rank-one weight subtraction ( $\mathbf{W}_{\text{new}} = \mathbf{W} - \alpha \cdot (\mathbf{r} \cdot \mathbf{r}^T) \mathbf{W}$ ) invariably alters the overall magnitude (norm) of the edited weights.<sup>8</sup> By shrinking the magnitude of specific pathways, standard ablation throws subsequent normalization layers out of equilibrium, causing the model to misinterpret perfectly valid inputs and resulting in downstream reasoning hallucinations.<sup>8</sup>

Norm-Preserving Biprojected Abliteration directly resolves this mechanical flaw. Rather than executing a blunt subtraction, this technique decomposes the weight matrix, modifies *only* the directional component, and then meticulously recombines the tensor to guarantee the

preservation of its original magnitude.<sup>8</sup>

## 6.1 Algorithmic Execution of Norm-Preservation

The protocol for ensuring norm preservation during an ablation edit proceeds through four rigorous mathematical steps<sup>8</sup>:

1. **Vector Normalization:** The purified target direction  $\mathbf{r}$  (ideally extracted via SRA) is normalized to unit length to isolate its pure directional influence:  $\hat{\mathbf{r}} = \frac{\mathbf{r}}{\|\mathbf{r}\|_2}$ .<sup>8</sup>
2. **Matrix Decomposition:** For every single row  $i$  of the targeted weight matrix  $\mathbf{W}$ , the specific magnitude  $m_i = \|\mathbf{W}_{i,:}\|_2$  is calculated and extracted. The entire matrix is subsequently decomposed into a diagonal matrix of scalar magnitudes  $\mathbf{M}$  and a purely directional matrix  $\hat{\mathbf{W}} = \mathbf{M}^{-1}\mathbf{W}$ .<sup>8</sup>
3. **Targeted Orthogonalization:** The projection coefficients—representing the specific alignment of each input dimension with the target direction—are computed as  $\mathbf{p} = \hat{\mathbf{r}}^T \hat{\mathbf{W}}$ . The target component is then surgically removed via a rank-1 update executed *exclusively* on the directional matrix, leaving magnitudes untouched<sup>8</sup>:

$$\hat{\mathbf{W}}_{\text{ablated}} = \hat{\mathbf{W}} - \alpha \cdot \hat{\mathbf{r}}\mathbf{p}^T$$

4. **Renormalization and Recombination:** Because the subtraction in step 3 altered the length of the vectors in the directional matrix, each row of  $\hat{\mathbf{W}}_{\text{ablated}}$  must be immediately re-normalized to unit length ( $\text{dim} = 1$ ) to form  $\hat{\mathbf{W}}_{\text{new}}$ . Finally, the original scalar magnitudes stored in matrix  $\mathbf{M}$  are restored<sup>8</sup>:

$$\mathbf{W}_{\text{new}} = \mathbf{M}\hat{\mathbf{W}}_{\text{new}}$$

This mathematical constraint serves as an absolute guarantee that  $\|\mathbf{W}_{\text{new},i,:}\|_2 = \|\mathbf{W}_{i,:}\|_2$  for all rows  $i$ . The learned importance structure and energy distribution of the matrix are perfectly preserved, while the computational flow is safely and permanently redirected away from the targeted behavior.<sup>8</sup> Empirical tests demonstrate that norm-preserving ablation not only prevents capability degradation but can actually *improve* performance on complex reasoning benchmarks compared to the baseline model.<sup>8</sup> By efficiently refocusing the model's capacity without destabilizing its internal variances, the model operates with enhanced clarity.<sup>8</sup>

## 7. Architectural Ablation and Differentiable

## Architecture Search (DARTS-GT)

While Representation Engineering (RepE) focuses on modulating latent vectors within pre-established, frozen architectures, highly advanced analytical operations frequently require structural ablation techniques to optimize the physical architecture of the model itself.<sup>38</sup> This is particularly critical when deploying models to resource-constrained environments where maximum computational efficiency must be achieved without sacrificing predictive accuracy, or when analyzing specialized architectures such as Graph Transformers (GTs).<sup>22</sup>

Modern high-end models increasingly rely on complex, heterogeneous attention mechanisms, moving away from homogeneous blocks applied uniformly across all depths.<sup>1</sup> In the context of Graph Transformers, ablation studies have been systematically operationalized to decouple structural encoding from feature representation.<sup>39</sup> The DARTS-GT (Differentiable Architecture Search for Graph Transformers) framework represents the bleeding edge of this methodology.<sup>39</sup>

DARTS-GT applies automated architecture search to select optimal Graph Neural Network (GNN) operators at specific, distinct layers, enabling depth-wise heterogeneity inside the transformer attention block itself.<sup>39</sup> However, architectures generated via automated search often become opaque black boxes, making it impossible to distinguish between meaningful structural patterns and spurious correlations.<sup>40</sup> To verify these dynamically generated designs without contamination, rigorous causal ablation is employed using advanced, instance-level metrics<sup>39</sup>:

- **Head-deviation:** This metric quantitatively measures the exact prediction changes that occur when specific attention heads are masked or ablated. It establishes a definitive baseline of causal load for every head in the network.<sup>39</sup>
- **Specialization:** Derived from head-deviation, specialization captures the dispersion of impacts across the network. A high specialization score indicates that only a very few, highly specific heads carry the causal load (and are thus absolutely critical to preserve). A low score indicates that the heads are acting redundantly and interchangeably, marking them as prime candidates for safe ablation and pruning.<sup>39</sup>
- **Focus:** This metric measures the consensus over specific graph regions, analyzing whether the top-performing heads agree on the most critical nodes.<sup>39</sup>

By systematically quantifying exactly which heads and nodes drive predictions via these ablation-grounded metrics, researchers can definitively map the architectural hierarchy. This allows for the aggressive pruning of non-essential components—drastically reducing parameter count and memory footprint—without inducing systemic failure or contaminating the model's specialized operational capabilities.<sup>39</sup>

## 8. Empirical Verification: Benchmark Profiling and the

# Ability Impact Score (AIS)

A central, systemic vulnerability in modern LLM engineering is the over-reliance on generalized, static benchmarks (e.g., MMLU, HellaSwag, GSM8K, ARC) to measure and validate model capability.<sup>9</sup> Benchmarks are notoriously ambiguous; achieving a high accuracy score on a dataset labeled "reasoning" or "commonsense" is routinely and falsely taken as absolute proof that the model actually possesses that generalized cognitive trait.<sup>19</sup>

In reality, LLMs frequently exploit superficial dataset artifacts, memorize specific evaluation patterns, or leverage an entangled mixture of entirely unrelated skills to achieve high scores without possessing any genuine underlying understanding.<sup>9</sup> This phenomenon creates the heavily documented "Performance-Perception Paradox," where models completely dominate public leaderboards but exhibit severe, unexpected capability gaps when deployed in real-world applications.<sup>9</sup>

When executing high-fidelity surgical model ablation (such as SRA or Norm-Preserving Biprojection), relying on traditional aggregate benchmark scores is entirely insufficient for proving that cross-distribution contamination has been avoided.<sup>43</sup> A model might maintain its overall score on a reasoning benchmark by compensating with superficial syntactic skills, masking the fact that its core logical circuits were severely damaged during the ablation.<sup>9</sup> To quantitatively prove that a surgical ablation successfully neutralized a target vector with absolute zero contamination of adjacent logic pathways, the analytical framework must employ **Benchmark Profiling** utilizing the **Ability Impact Score (AIS)**.<sup>9</sup>

## 8.1 The Architecture of the Ability Impact Score (AIS)

Benchmark Profiling is a highly systematic methodology designed to diagnose the exact functional abilities required by LLM benchmarks.<sup>9</sup> It decomposes overall performance into cognitively grounded metrics, combining gradient-based importance scoring with targeted, localized parameter ablation.<sup>9</sup>

The execution of Benchmark Profiling follows a rigorous four-step pipeline<sup>9</sup>:

1. **Cognitive Operationalization:** Ten fundamental, distinct cognitive abilities are strictly defined (e.g., Deductive Reasoning, Contextual Recall, Quantitative Reasoning, Analogical Reasoning, Spatial Reasoning) based on established, peer-reviewed models of human intelligence.<sup>9</sup>
2. **Gradient-Based Importance Scoring:** Utilizing a dedicated, isolated diagnostic dataset for each specific ability, the model's internal parameters are ranked by gradient-based importance. This precisely quantifies the sensitivity of the model's loss function to specific perturbations associated *only* with that distinct cognitive skill.<sup>9</sup>
3. **Targeted Parameter Ablation Constraint:** To isolate the capability, the top **1.024%** of Multi-Layer Perceptron (MLP) parameters explicitly associated with that specific ability



are temporarily zeroed out.<sup>9</sup> Extensive preliminary sweeps dictate that **1.024%** represents the optimal, minimal ablation budget capable of producing a clear, ability-specific signal degradation without inflicting any unnecessary collateral damage on the rest of the network.<sup>9</sup>

4. **AIS Computation:** Both the baseline model and the newly ability-ablated model are executed on the downstream benchmarks. The resulting degradation in task accuracy yields the Ability Impact Score (AIS).<sup>9</sup> The AIS mathematically quantifies exactly how strongly a given benchmark depends on a specific, isolated cognitive ability.<sup>9</sup>

The formulation for the Ability Impact Score is represented as <sup>44</sup>:

$$AIS_b^{(a)} = \frac{P_b(\Theta) - P_b(\Theta^{(a)})}{P_b(\Theta) - P_b^{chance}}$$

(Where  $P_b(\Theta)$  signifies the baseline performance of the model,  $P_b(\Theta^{(a)})$  signifies the performance of the ablated model, and  $P_b^{chance}$  represents the statistical probability of random chance <sup>44</sup>).

## 8.2 Analytical Application of AIS Post-Intervention

Following the execution of a behavioral ablation protocol, the modified foundational model must be immediately subjected to the exhaustive Benchmark Profiling suite.<sup>9</sup>

If the surgical ablation was successfully contained (i.e., the target behavior was removed without causing "Ghost Noise"), the AIS radar plots for the modified model will match the baseline model's plots identically across all protected domains.<sup>19</sup> For instance, the AIS values for Quantitative Reasoning and Inductive Logic should display absolute zero deviation.<sup>19</sup> Any statistically significant deviation in the AIS profile acts as a highly sensitive diagnostic alarm, definitively indicating that the target vector was not sufficiently orthogonalized and that the underlying MLP parameters have been contaminated.<sup>2</sup>

Furthermore, Benchmark Profiling provides profound structural insights into modern foundational models like Llama-3.1-8B-Instruct.<sup>9</sup> Analyses utilizing the AIS framework reveal that code-generation benchmarks (such as HumanEval and MBPP) yield massive AIS values across nearly every defined cognitive axis, forming an almost complete geometric disk on diagnostic radar plots.<sup>9</sup> This explicitly demonstrates that coding proficiency relies on highly entangled, multi-skill dependencies.<sup>9</sup> Conversely, certain abilities register *negative* AIS values (e.g., ablating Long-Term Knowledge actually increases LogiQA accuracy by 1-2 percentage points), highlighting paradoxical instances where models rely on spurious cues that actively conflict with their core reasoning tasks.<sup>19</sup> Understanding this highly complex, interconnected topological map is absolutely paramount prior to attempting any form of latent vector intervention, ensuring the researcher is fully aware of the precise cognitive dependencies at



risk during ablation.

| Benchmark Evaluation Focus         | Cognitive Dependency Complexity | AIS Mapping Insight                                                                               | Post-Ablation Vulnerability                                                                                  |
|------------------------------------|---------------------------------|---------------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------|
| Code Generation (MBPP, HumanEval)  | Extremely High                  | Draws heavily on all 10 operationalized abilities. Forms complete geometric disks on radar plots. | Highly susceptible to collateral damage. Any misaligned ablation will instantly degrade coding capabilities. |
| Question Answering (CommonsenseQA) | Moderate to High                | Peaks sharply in Inductive Reasoning, with secondary reliance on Deductive and Spatial reasoning. | Vulnerable to structural orthogonalization errors that impact inductive logic pathways.                      |
| Logical Reasoning (LogiQA)         | Complex / Paradoxical           | Exhibits negative AIS values for Long-Term Knowledge, indicating reliance on spurious artifacts.  | Paradoxically, careless ablation might <i>improve</i> scores if it destroys distracting knowledge artifacts. |
| Temporal / Analogical (WinoGrande) | Narrow                          | Narrowest profile, but still relies on multiple distinct competencies.                            | Generally stable, but requires verification that analogical pathways remain undisturbed by style confounds.  |

## 9. Strategic Conclusions and Operational

# Implementation

The manipulation of high-end, multi-billion parameter architectures requires a fundamental departure from legacy interpretability frameworks. The transition from naive, zero-value ablation to high-fidelity Representation Engineering necessitates a profound mathematical respect for the geometric complexity and deep polysemantic nature of neural networks.<sup>2</sup>

The empirical data clearly and consistently demonstrates that highly targeted behaviors—including strict safety limits, complex refusal mechanisms, and stylistic constraints—can be permanently modified or entirely eradicated from foundational models without ever relying on resource-intensive Supervised Fine-Tuning (SFT) or Direct Preference Optimization (DPO).<sup>1</sup> However, to execute this level of surgical modification without inducing widespread systemic contamination, analytical teams must adhere strictly to the following operational parameters:

1. **Mandatory Spectral Residualization:** Researchers must never perform rank-one weight subtraction using a raw, "dirty" vector calculated solely from mean activation differences. The vector must always be passed through a rigorously curated Concept Atom Registry (utilizing independent, data-light prompt sets) to execute ridge-regularized regression. This is the only proven method to protect essential "Shield" capabilities from unintended excision.<sup>2</sup>
2. **Enforce Norm-Preservation at all Layers:** Direct weight edits must be mathematically decomposed. Ablation must occur strictly on the directional matrix, followed by aggressive renormalization to unit length prior to recombining with the original scalar magnitudes. This protocol is non-negotiable for preserving the stability of existing layer normalization and preventing downstream reasoning hallucinations.<sup>8</sup>
3. **Implement Peak Ablation for Non-Destructive Testing:** When conducting exploratory analysis, hypothesis testing, or non-permanent inference-time steering, researchers should utilize modal "Peak Activation Centering" rather than zero or mean ablation. This minimizes the injection of out-of-distribution mathematical noise into the residual stream, ensuring a cleaner analytical read.<sup>3</sup>
4. **Continuous AIS Diagnostics as Integrity Gates:** Leverage comprehensive Benchmark Profiling and the Ability Impact Score (AIS) to continuously audit the semantic integrity of the model. Traditional accuracy metrics are blind to underlying geometric degradation; AIS provides a definitive, granular, and mathematically rigorous audit of cognitive preservation.<sup>9</sup>

By strictly enforcing these mathematical algorithms and procedural constraints, advanced analytical teams can successfully isolate, uniquely modulate, or permanently ablate any targeted behavioral phenomenon within a foundational model. This methodology achieves absolute precision while guaranteeing the uncompromising preservation of the model's core computational integrity and generalized intelligence.

## Works cited

1. LLaMA-3 8B-Instruct Overview - Emergent Mind, accessed on March 19, 2026, <https://www.emergentmind.com/topics/llama-3-8b-instruct>
2. Surgical Refusal Ablation: Disentangling Safety from Intelligence via Concept-Guided Spectral Cleaning - arXiv, accessed on March 19, 2026, <https://arxiv.org/html/2601.08489v1>
3. Investigating Neuron Ablation in Attention Heads: The Case for Peak Activation Centering, accessed on March 19, 2026, <https://arxiv.org/html/2408.17322v1>
4. Surgical Refusal Ablation: Disentangling Safety from Intelligence via Concept-Guided Spectral Cleaning - arXiv, accessed on March 19, 2026, <https://arxiv.org/pdf/2601.08489>
5. Surgical Refusal Ablation: Disentangling Safety from Intelligence via Concept-Guided Spectral Cleaning - ResearchGate, accessed on March 19, 2026, [https://www.researchgate.net/publication/399754691\\_Surgical\\_Refusal\\_Ablation\\_Disentangling\\_Safety\\_from\\_Intelligence\\_via\\_Concept-Guided\\_Spectral\\_Cleaning](https://www.researchgate.net/publication/399754691_Surgical_Refusal_Ablation_Disentangling_Safety_from_Intelligence_via_Concept-Guided_Spectral_Cleaning)
6. Abliteration - Learn AI, accessed on March 19, 2026, <https://ai.miraheze.org/wiki/Abliteration>
7. Refusal in LLMs is mediated by a single direction - LessWrong, accessed on March 19, 2026, <https://www.lesswrong.com/posts/jGuXSZgv6qfdhMCuJ/refusal-in-llms-is-mediated-by-a-single-direction>
8. Norm-Preserving Biprojected Abliteration - Hugging Face, accessed on March 19, 2026, <https://huggingface.co/blog/grimjim/norm-preserving-biprojected-abliteration>
9. Benchmark Profiling: Mechanistic Diagnosis of LLM Benchmarks - arXiv, accessed on March 19, 2026, <https://arxiv.org/html/2510.01232v1>
10. An Introduction to Representation Engineering - an activation-based ..., accessed on March 19, 2026, <https://www.alignmentforum.org/posts/3ghj8EuKzwd3MQR5G/an-introduction-to-representation-engineering-an-activation>
11. Representation Engineering for Large-Language Models: Survey and Research Challenges, accessed on March 19, 2026, <https://arxiv.org/html/2502.17601v1>
12. Refusal in Language Models Is Mediated by a Single Direction - OpenReview, accessed on March 19, 2026, [https://openreview.net/forum?id=pH3X AQME6c&referrer=%5Bthe%20profile%20of%20Wes%20Gurnee%5D\(%2Fprofile%3Fid%3D~Wes\\_Gurnee1\)](https://openreview.net/forum?id=pH3X AQME6c&referrer=%5Bthe%20profile%20of%20Wes%20Gurnee%5D(%2Fprofile%3Fid%3D~Wes_Gurnee1))
13. jim-plus/llm-abliteration - GitHub, accessed on March 19, 2026, <https://github.com/jim-plus/llm-abliteration/>
14. GitHub - Orion-zhen/abliteration: Make ablated models with transformers, easy and fast, accessed on March 19, 2026, <https://github.com/Orion-zhen/abliteration>
15. NousResearch/llm-abliteration: Make ablated models ... - GitHub, accessed on March 19, 2026, <https://github.com/NousResearch/llm-abliteration>
16. 1 Introduction - arXiv, accessed on March 19, 2026,

- <https://arxiv.org/html/2512.13655v2>
17. Controlled Ablation Study: Methods & Applications, accessed on March 19, 2026, <https://www.emergentmind.com/topics/controlled-ablation-study>
  18. Refusal in Language Models Is Mediated by a Single Direction - ResearchGate, accessed on March 19, 2026, [https://www.researchgate.net/publication/397221072\\_Refusal\\_in\\_Language\\_Models\\_Is\\_Mediated\\_by\\_a\\_Single\\_Direction](https://www.researchgate.net/publication/397221072_Refusal_in_Language_Models_Is_Mediated_by_a_Single_Direction)
  19. Ability Impact Score radar plots for the ten benchmarks profiled on... - ResearchGate, accessed on March 19, 2026, [https://www.researchgate.net/figure/Ability-Impact-Score-radar-plots-for-the-ten-benchmarks-profiled-on\\_fig2\\_396143515](https://www.researchgate.net/figure/Ability-Impact-Score-radar-plots-for-the-ten-benchmarks-profiled-on_fig2_396143515)
  20. Understanding Ablation Studies for Product Teams, accessed on March 19, 2026, <https://www.productteacher.com/quick-product-tips/ablation-studies-for-product-teams>
  21. Ablation (artificial intelligence) - Wikipedia, accessed on March 19, 2026, [https://en.wikipedia.org/wiki/Ablation\\_\(artificial\\_intelligence\)](https://en.wikipedia.org/wiki/Ablation_(artificial_intelligence))
  22. Ablation - Deepgram, accessed on March 19, 2026, <https://deepgram.com/ai-glossary/ablation>
  23. Ablation Study - Tasq.ai, accessed on March 19, 2026, <https://www.tasq.ai/glossary/ablation-study/>
  24. Investigating Neuron Ablation in Attention Heads: The Case for Peak Activation Centering - CEUR-WS.org, accessed on March 19, 2026, [https://ceur-ws.org/Vol-3793/paper\\_21.pdf](https://ceur-ws.org/Vol-3793/paper_21.pdf)
  25. Taxonomy, Opportunities, and Challenges of Representation Engineering for Large Language Models - arXiv, accessed on March 19, 2026, <https://arxiv.org/html/2502.19649v5>
  26. [Quick Review] Refusal in Language Models Is Mediated by a Single Direction - Liner, accessed on March 19, 2026, <https://liner.com/review/refusal-in-language-models-is-mediated-by-a-single-direction>
  27. Interpreting Transformers Through Attention Head Intervention - arXiv, accessed on March 19, 2026, <https://arxiv.org/html/2601.04398v4>
  28. Emergent Modularity in Pre-trained Transformers | Request PDF - ResearchGate, accessed on March 19, 2026, [https://www.researchgate.net/publication/372917801\\_Emergent\\_Modularity\\_in\\_Pre-trained\\_Transformers](https://www.researchgate.net/publication/372917801_Emergent_Modularity_in_Pre-trained_Transformers)
  29. Optimal ablation for interpretability | Lucas Janson, accessed on March 19, 2026, [https://lucasjanson.fas.harvard.edu/papers/Optimal\\_Ablation\\_For\\_Interpretability-Li\\_Janson-2024.pdf](https://lucasjanson.fas.harvard.edu/papers/Optimal_Ablation_For_Interpretability-Li_Janson-2024.pdf)
  30. Representation Engineering: A Top-Down Approach to AI Transparency - arXiv, accessed on March 19, 2026, <https://arxiv.org/html/2310.01405v4>
  31. Rethinking The Reliability of Representation Engineering in Large Language Models, accessed on March 19, 2026, <https://openreview.net/forum?id=sYJQEgkkaI>
  32. Refusal Direction Surgery in Transformers - Emergent Mind, accessed on March

- 19, 2026, <https://www.emergentmind.com/topics/refusal-direction-surgery>
33. DeepRefusal: Probabilistic Refusal Ablation - Emergent Mind, accessed on March 19, 2026, <https://www.emergentmind.com/topics/probabilistic-refusal-direction-ablation-deeprefusal>
34. Refusal in Language Models Is Mediated by a Single Direction - NIPS, accessed on March 19, 2026, [https://proceedings.neurips.cc/paper\\_files/paper/2024/file/f545448535dfde4f9786555403ab7c49-Paper-Conference.pdf](https://proceedings.neurips.cc/paper_files/paper/2024/file/f545448535dfde4f9786555403ab7c49-Paper-Conference.pdf)
35. Uncensor any LLM with ablation - Hugging Face, accessed on March 19, 2026, <https://huggingface.co/blog/mlabonne/ablation>
36. Refusal in Language Models Is Mediated by a Single Direction - OpenReview, accessed on March 19, 2026, <https://openreview.net/pdf?id=EqF16oDVff>
37. Null It Out: Guarding Protected Attributes by Iterative Nullspace Projection - ResearchGate, accessed on March 19, 2026, [https://www.researchgate.net/publication/343299041\\_Null\\_It\\_Out\\_Guarding\\_Protected\\_Attributes\\_by\\_Iterative\\_Nullspace\\_Projection](https://www.researchgate.net/publication/343299041_Null_It_Out_Guarding_Protected_Attributes_by_Iterative_Nullspace_Projection)
38. Utilizing Large Language Models for Ablation Studies in Machine Learning and Deep Learning - Diva-Portal.org, accessed on March 19, 2026, <https://www.diva-portal.org/smash/get/diva2:1941572/FULLTEXT01.pdf>
39. DARTS-GT: Differentiable Architecture Search for Graph Transformers with Quantifiable Instance-Specific Interpretability Analysis - arXiv, accessed on March 19, 2026, <https://arxiv.org/pdf/2510.14336>
40. (PDF) DARTS-GT: Differentiable Architecture Search for Graph Transformers with Quantifiable Instance-Specific Interpretability Analysis - ResearchGate, accessed on March 19, 2026, [https://www.researchgate.net/publication/396542232\\_DARTS-GT\\_Differentiable\\_Architecture\\_Search\\_for\\_Graph\\_Transformers\\_with\\_Quantifiable\\_Instance-Specific\\_Interpretability\\_Analysis](https://www.researchgate.net/publication/396542232_DARTS-GT_Differentiable_Architecture_Search_for_Graph_Transformers_with_Quantifiable_Instance-Specific_Interpretability_Analysis)
41. Benchmark Profiling: Mechanistic Diagnosis of LLM Benchmarks - ACL Anthology, accessed on March 19, 2026, <https://aclanthology.org/2025.emnlp-main.789.pdf>
42. Gyuho Shim's research works - ResearchGate, accessed on March 19, 2026, <https://www.researchgate.net/scientific-contributions/Gyuho-Shim-2320965231>
43. A Complete Guide to LLM Benchmark Categories - Galileo AI, accessed on March 19, 2026, <https://galileo.ai/blog/llm-benchmarks-categories>
44. HellaSwag Benchmark - Emergent Mind, accessed on March 19, 2026, <https://www.emergentmind.com/topics/hellaswag-benchmark>